

Exercise sheet 2

Prepare the below such that you are able to discuss it on Wednesday 10th March

Exercise 1 Continuous time stochastic SIR model

The simple continuous time stochastic SIR model was in the lecture formulated as a stochastic process with piecewise constant conditional intensity function (CIF).

- a) Write an R function which creates one realization of this model with inputs $\theta = (\beta, \gamma)'$ and $(s(0), i(0))'$. The function should return a matrix with 3 columns where each row contains the event time and the number of susceptibles at that event time.

Solution:

```
# Function for converting missings to infinite values
NA2Inf <- function(x) {x[is.na(x)] <- Inf; return(x)}

simulate_sir <- function(theta, si0) {
  beta <- theta[1]
  gamma <- theta[2]
  t <- 0

  # Results should be given in a three-dimensional matrix
  # with columns t, S(t), I(t)
  res <- matrix(c(t, si0), ncol = 3)
  colnames(res) <- c("t", "s", "i")

  # If I(t) > 0 continue
  while(tail(res, n = 1)[, 3] > 0) { # Check for above requirement
    sit <- tail(res, n = 1)[, 2 : 3]
    # Intensity of the two competing events F
    cif <- c("S->I" = beta * sit[1] * sit[2],
            "I->R" = gamma * sit[2]) # Transitions
    # Simulate waiting times from an exponential distribution and
    # apply the function that makes missing infinite
    waiting_time <- NA2Inf(rexp(2, rate = cif))

    if (waiting_time[1] < waiting_time[2]) { #S->I event occurs
      t <- t + waiting_time[1]
      sit <- c(sit[1] - 1, sit[2] + 1)
    } else {
      t <- t + waiting_time[2]
      sit <- c(sit[1], sit[2] - 1)
    }
    res <- rbind(res, c(t, sit))
  }
  return(as.data.frame(res))
}
```


- b) Test your function from the previous question with the inputs values $\theta = (0.1, 0.3)'$ and

$(s(0), i(0))' = (20, 1)'$.

(Hint: A graphical representation may help visualise the results)

Solution:

```
# Use the inputs provided
set.seed(20210303)
traj <- simulate_sir(theta = c(beta = 0.1, gamma = 0.3), si0 = c(20, 1))
# Plot the results
par(mfrow = c(1, 2))
plot(traj$t, traj$s, type = "s", ylab = "S(t)", xlab = "t")
abline(v = traj$t[which(traj$s == 0)], col = 4, lty = 2)
# Add line for first instance where S(t) = 0
plot(traj$t, traj$i, type = "s", ylab = "I(t)", xlab = "t")
abline(v = traj$t[which(traj$s == 0)], col = 4, lty = 2)
```



```
# Add line for first instance where S(t) = 0
```

The curves for changes to the susceptible and infected compartments look as would be expected from a simple SIR model (Susceptibles are decreasing S-shape and I is bell-shaped). The peak of I is around the time where S has depleted and I continues to decrease thereafter (this makes sense as there are no more people available to be infected). The curve for the R component would similarly be expected to grow from around this time point.

Exercise 2 Reed-Frost model

The dynamics of the Reed-Frost model can be described by the discrete-time Markov chain

$$Y_{t+1}|x_t, y_t \sim \text{Bin}(x_t, 1 - (1 - w)^{y_t}),$$
$$X_{t+1} = X_t - Y_{t+1},$$

with $x_0 = n$ susceptible and $y_0 = m$ infectious individuals initially and w is the probability of an infectious contact between susceptible individual i and infectious individual j in interval $(t, t+1]$. The Reed-Frost model relates to the Kermack-McKendrick SIR formulation of models since Y describes changes to the I compartment and X describes changes to the S compartment.

- a) Create an R function which takes inputs w , x_0 and y_0 which simulates one instance of a Reed-Frost model with probability w and initial values $x_0 = x_0$ and $y_0 = y_0$. The function should return a **matrix** with two columns such that each row contains (x_t, y_t) for $t = 0, 1, 2, \dots$ (*Hint: You may want to use `rbind`*)

Solution:

```
reed_frost <- function(w, x0, y0) {
  XY <- matrix(c(x0, y0), ncol = 2) # Initial values
  t <- 1 # Initialise step in loop
  while (all(XY[t, ] > 0)) { # Stopping criterion for while-loop
    theta <- 1 - (1 - w) ^ XY[t, 2] # Parameter for binomial distribution
    ytp1 <- rbinom(n = 1, size = XY[t, 1], prob = theta) # Calculate  $Y_{t+1}$ 
    xtp1 <- XY[t, 1] - ytp1 # Calculate  $X_{t+1}$  with  $Y_{t+1}$  from above
    XY <- rbind(XY, c(xtp1, ytp1)) # Add to matrix object
    t <- t + 1 # Increase step
  }
  return(XY)
}
```

- b) Test this function with the probability $w = 0.15$ and initialisation $x_0 = 20, y_0 = 1$. (*Hint: You may want to set a seed such that you can use the output elsewhere*)

Solution:

```
set.seed(20210224)
reed_frost(w = 0.15, x0 = 20, y0 = 1)

##      [,1] [,2]
## [1,]   20   1
## [2,]   18   2
## [3,]   11   7
## [4,]    5   6
## [5,]    3   2
## [6,]    2   1
## [7,]    2   0
```

- c) Recall the final size of the epidemic is $Z = Y_1 + Y_2 + Y_3 + \dots$. Given the output of your function from question 1 a), construct another R function which takes the output of the previous function as input and calculates Z . Calculate this for the values given in question 1 b).

Solution:

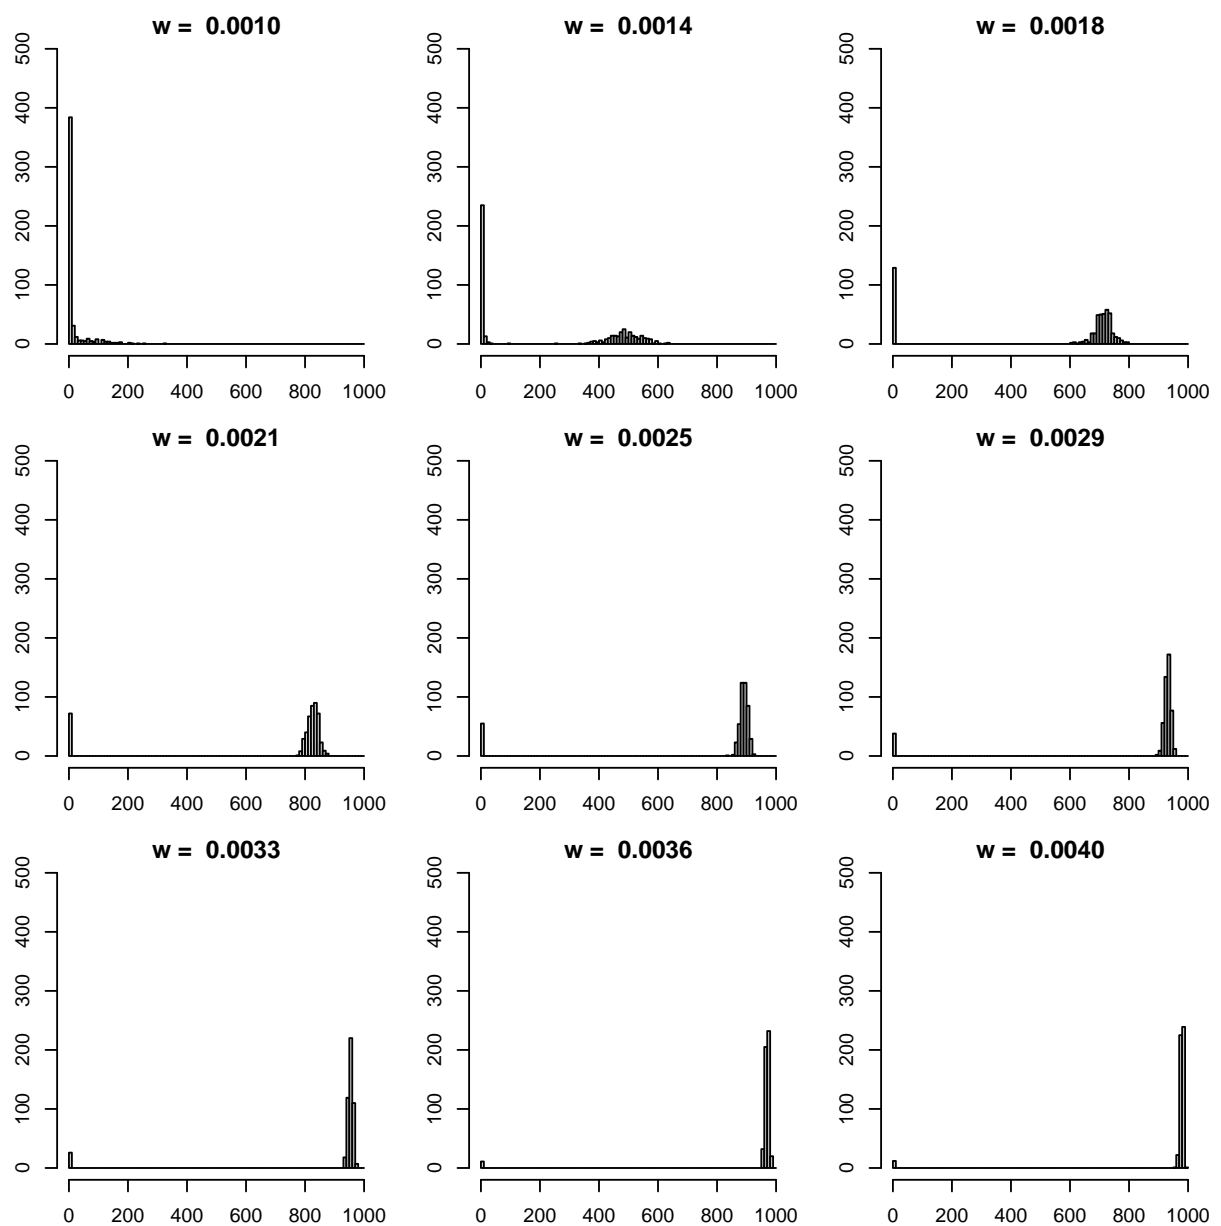
```
final_size <- function(mat) {
  return(sum(mat[- 1, # To ensure summing from  $Y_1$  not  $Y_0$ 
                2, # Interested in second column (the  $Y_i$ 's)
                drop = FALSE]))
}
set.seed(20210224)
final_size(reed_frost(w = 0.15, x0 = 20, y0 = 1))

## [1] 18
```

- d) Now consider $x_0 = 1000$ and $y = 1$. Describe how the final size Z relates to w . (*Hint: A graphical representation may help visualise what is going on*)

Solution:

```
x0 <- 1000
par(mfrow = c(3, 3), mar = c(3, 3, 1, 1))
for (w in seq(0.001, 0.004, length = 9)) {
  hist(
    replicate(500, final_size(reed_frost(w = w, x0 = x0, y0 = 1))),
    # Calculate 500 instances of the Reed-Frost model with
    # x = 1000, y = 1
    # for various w
    breaks = seq(0, x0, by = 10),
    main = paste("w = ", sprintf("%.4f", w)), # Add descriptive titles
    xlab = "Z",
    freq = TRUE, ylim = c(0, 500))
}
```



We see that when w is smaller, i.e. the probability of an infectious contact decreases, the final size of the epidemic also decreases.

e) Given the values

	t_0	t_1	t_2	t_3	t_4	t_5
x_t	20	18	12	4	2	0
y_t	1	2	6	8	2	2

estimate w and its associated 95% Wald confidence interval using maximum likelihood estimation.

(*Hint:* You may want to use a parameter transformation for the probability parameter, e.g. based on the inverse cdf of the logistic distribution as well as `optim`)

Solution: Recall that the Wald CI is given by α

```
l <- function(w_logit, x, y) {  
  if (length(x) != length(y)) { stop("x and y need to be the same length") }  
  K <- length(x)  
  w <- plogis(w_logit)  
  theta <- 1 - (1 - w) ^ y  
  #Compute loglik  
  return(sum(dbinom(y[- 1], size = x[- K], prob = theta[- K], log = TRUE)))  
}  
mle <- optim(par = 0, fn = l, method = "BFGS",  
            x = c(20, 18, 12, 4, 2, 0),  
            y = c(1, 2, 6, 8, 2, 2),  
            control = list(fnscale = - 1), hessian = TRUE)  
#Maximum likelihood estimator  
plogis(mle$par)  
  
## [1] 0.1574987  
  
#95% confidence interval  
c(plogis(mle$par - qnorm(0.975) * sqrt(- 1 / mle$hess)),  
  plogis(mle$par + qnorm(0.975) * sqrt(- 1 / mle$hess)))  
  
## [1] 0.1027789 0.2337612
```